
Efficient KV Cache Management for Long-Context LLM Inference: Lazy Pruning, Slack, and Staged Eviction

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Long-context LLM inference is bottlenecked by KV-cache growth and per-token
2 runtime overhead. StreamingLLM [10] bounds KV-cache size by retaining attention sinks and a recent-token window, but practical decoding can suffer from (i)
3 non-trivial KV pruning overhead and (ii) quality degradation caused by periodic
4 hard evictions, especially on long-form generation. We propose three lightweight,
5 training-free KV-cache management mechanisms on top of StreamingLLM: **Lazy**
6 **Pruning** (amortize pruning by triggering after R tokens of overflow), **Slack**
7 (a hard-cap buffer that enables staged evictions under bounded memory), and
8 **Max_Drop** (limit the maximum tokens evicted per prune event). We evaluate
9 on Pythia-2.8B using WikiText-103 (sanity) and PG19 (long-context), and report
10 TTFT/TPOT/throughput/memory/perplexity trade-offs. We also document why
11 FlashAttention, speculative decoding, quantization, and CUDA operator fusion did
12 not improve batch-1 streaming decode under our setting.
13

1 Introduction

15 Long-context LLM inference is challenging because both computation and memory scale with the
16 context length. During autoregressive decoding, the KV cache grows linearly with the number of
17 processed tokens, increasing attention cost and memory traffic. StreamingLLM [10] mitigates this
18 by keeping only (i) the first S “sink” tokens and (ii) a sliding window of the most recent W tokens,
19 thereby bounding the effective KV-cache length.

20 **Motivation.** In our setting (Pythia-2.8B on NVIDIA A800), once attention length is bounded,
21 decode becomes dominated by MLP and framework/launch overhead. Meanwhile, KV pruning itself
22 can incur additional overhead due to slicing/copy and rotary-position re-alignment. More importantly,
23 periodic hard eviction introduces distribution shifts that manifest as token-level NLL spikes and
24 perplexity (PPL) degradation on long-form datasets. These observations motivate KV-management
25 mechanisms that (i) reduce amortized pruning overhead and (ii) stage evictions to stabilize quality,
26 without modifying attention kernels or model weights.

27 **Contributions.** We propose three training-free, modular mechanisms on top of StreamingLLM: (i)
28 **Lazy Pruning**, which triggers pruning only after the cache overflow reaches R tokens; (ii) **Slack**,
29 a hard capacity buffer of σ tokens that supports staged pruning; (iii) **Max_Drop**, which limits the
30 maximum tokens evicted per prune event to δ . We provide a reproducible evaluation pipeline and
31 summarize negative results for methods that are effective in other regimes but do not improve batch-1
32 streaming decode.

33 **2 Related Work**

34 **Long-context inference and KV-cache compression.** Test-time KV-cache management is a widely
 35 used approach to make long-context inference feasible, including retention/eviction heuristics and
 36 cache compression [10, 11, 8, 6]. Our work builds on StreamingLLM’s Start+Recent rule and focuses
 37 on *training-free* extensions that target the practical costs of pruning and eviction-induced distribution
 38 shifts. **Kernel- and decoding-level accelerations.** Attention kernels such as FlashAttention aim
 39 to reduce attention IO cost [3, 2], while speculative decoding accelerates generation by drafting
 40 tokens [7, 1]. Post-training quantization reduces compute/memory but its benefits depend on hardware
 41 and runtime regime [4, 5]. We empirically find these do not improve batch-1 streaming decode in our
 42 setting and analyze the underlying reasons (Sec. 5).

43 **3 Method**

44 We first define Start+Recent streaming, then formalize our three mechanisms.

45 **3.1 Preliminaries: Start+Recent StreamingLLM**

46 Let $x_{1:t}$ be the processed prefix at decode step t . For each layer ℓ , let $(\mathbf{K}_t^\ell, \mathbf{V}_t^\ell)$ denote the KV cache
 47 with sequence length L_t along the cache dimension. Start+Recent streaming retains (i) the first S
 48 sink tokens, and (ii) the most recent W tokens. Define the soft capacity

$$C_0 \triangleq S + W. \quad (1)$$

49 Our code also supports optional *non-core* retained tokens (e.g., overlap/refresh); we denote their total
 50 budget by $B \geq 0$ and define the effective soft capacity

$$C \triangleq C_0 + B. \quad (2)$$

51 Unless stated otherwise, we use $B = 0$ in paper experiments to isolate our three mechanisms. When
 52 pruning is performed, the retained index set is

$$\mathcal{I}_t = \{0, 1, \dots, S - 1\} \cup \{L_t - W, \dots, L_t - 1\}, \quad (3)$$

53 with the recent segment clamped to avoid overlap with the sink segment.

54 **RoPE consistency.** For rotary-position-embedding models (e.g., Pythia/GPT-NeoX), pruning
 55 changes token positions in the cache; therefore, cached keys must be re-aligned to the new po-
 56 sitions. Let $p \in \mathcal{I}_t$ be an old token position and $p' = f_t(p) \in \{0, \dots, |\mathcal{I}_t| - 1\}$ be its new position
 57 after compaction. RoPE encodes position via per-dimension frequencies ω_i (from `inv_freq`).
 58 Re-alignment is equivalent to applying a *delta-rotation* with $\Delta p = p' - p$:

$$\begin{bmatrix} k'_{2i} \\ k'_{2i+1} \end{bmatrix} = \begin{bmatrix} \cos(\omega_i \Delta p) & -\sin(\omega_i \Delta p) \\ \sin(\omega_i \Delta p) & \cos(\omega_i \Delta p) \end{bmatrix} \begin{bmatrix} k_{2i} \\ k_{2i+1} \end{bmatrix}. \quad (4)$$

59 Our implementation exploits that Start+Recent yields a constant Δp for the contiguous recent block (a
 60 constant shift), and falls back to the general per-token re-rotation when tokens are explicitly relocated
 61 (e.g., refresh tokens).

62 **3.2 Lazy Pruning**

63 Naïve streaming prunes as soon as the cache exceeds C , which can add overhead due to slicing/copy
 64 and KV relocation. We instead amortize pruning by allowing the cache to exceed C by up to $R - 1$
 65 tokens before triggering a prune. Define overflow:

$$\text{overflow}_t \triangleq L_t - C. \quad (5)$$

66 Let R be the *overflow allowance* hyperparameter (in code: `compress_every`). We allow $R = 0$ to
 67 disable pruning for debugging (unbounded memory), and use $R \geq 1$ in all bounded-memory settings.
 68 Lazy pruning triggers only when overflow reaches R :

$$\text{PruneTrigger}(t) \triangleq \mathbb{1}[R > 0] \cdot \mathbb{1}[\text{overflow}_t \geq R]. \quad (6)$$

69 This matches the implementation semantics and avoids the ambiguity that would arise from interpret-
 70 ing $R = 0$ as “prune immediately”.

Algorithm 1 KV pruning with Lazy Pruning, Slack, and Max_Drop

Require: sink size S , window size W , extra budget B (default 0), overflow allowance R (0 disables), slack σ , max_drop δ

- 1: $C_0 \leftarrow S + W, C \leftarrow C_0 + B, H \leftarrow C + \sigma$
- 2: observe cache length L_t
- 3: **if** $R = 0$ **then**
- 4: **return**
- 5: **end if**
- 6: **if** $L_t \leq C$ **then**
- 7: **return**
- 8: **end if**
- 9: **if** $L_t - C < R$ **then**
- 10: **return**
- 11: **end if**
- 12: **if** $\delta = 0$ **then**
- 13: $L_t^* \leftarrow C$
- 14: **else**
- 15: $L_t^* \leftarrow \min(\max(L_t - \delta, C), H)$
- 16: **end if**
- 17: keep first S tokens and most recent $(L_t^* - S)$ tokens
- 18: prune KV cache and re-align RoPE positions

71 **3.3 Slack (Hard-Cap Buffer)**72 Slack introduces a *hard* capacity that bounds temporary cache growth during staged pruning. Specifi-
73 cally, we define

$$H \triangleq C + \sigma, \quad (7)$$

74 where $\sigma \geq 0$ is a fixed slack budget. In our implementation, σ does *not* change the pruning trigger; it
75 caps the post-prune target length when combined with Max_Drop (Sec. 3.4).76 **3.4 Max_Drop (Staged Eviction)**77 Periodic pruning can cause “eviction cliffs” when a large block is removed at once, which we observed
78 as token-level NLL spikes on long-form generation. Max_Drop limits the maximum number of
79 tokens evicted in a single prune event.80 Let $\delta \geq 0$ be the maximum drop size. When pruning is triggered, we set a target retained length

$$L_t^* = \begin{cases} C, & \delta = 0 \\ \min(\max(L_t - \delta, C), H), & \delta > 0 \end{cases} \quad (8)$$

81 and prune the cache to keep exactly L_t^* tokens using the Start+Recent rule. This turns a single large
82 eviction into multiple smaller evictions across steps while remaining within the hard capacity H .83 **Hyperparameter interactions and a toy example.** When $\delta > 0$, staged eviction can leave
84 $L_t^* > C$, so pruning may trigger again soon; this is intentional. A practical rule-of-thumb to avoid
85 pruning *every* step when the cache saturates at H is to choose $\sigma < R$ (so that after pruning to H ,
86 the overflow remains $< R$ until enough new tokens arrive). For example, with $(S, W, R, \sigma, \delta) =$
87 $(4, 2044, 32, 16, 32)$ (so $C = 2048, H = 2064$), if a step reaches $L_t = 2090$ (overflow = 42), we
88 prune to $L_t^* = \min(\max(2090 - 32, 2048), 2064) = 2058$ and continue staging further evictions as
89 decoding proceeds.90 **3.5 Algorithm Summary**91 Algorithm 1 summarizes the pruning control logic (RoPE re-alignment and buffer management are
92 omitted for clarity).

Table 1: Main results on PG19 (long-context, *fill values via script output*).

Method	TPOT \downarrow	Speedup \uparrow	PPL \downarrow	Peak Mem (MB) \downarrow
Baseline (Full KV)	100.53	1.00 \times	19.761	5722
StreamingLLM (MIT)	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]
Ours (Lazy/Slack/Max_Drop)	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]

Table 2: Sanity-check results on WikiText-103 (short-context).

Method	TPOT \downarrow	Speedup \uparrow	PPL \downarrow
Baseline (Full KV)	98.86	1.00 \times	9.359
StreamingLLM (MIT)	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]
Ours (Lazy/Slack/Max_Drop)	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]

93 4 Experiments

94 4.1 Setup

95 We evaluate on Pythia-2.8B. We use WikiText-103 as a short-context sanity check and PG19 as
96 the primary long-context benchmark [9]. We report TTFT, TPOT (ms/token), throughput (tok/s),
97 total runtime, peak GPU memory, and perplexity (PPL). **Baselines and fairness.** Our long-context
98 evaluation must respect the model’s maximum position length. Therefore, we compare methods
99 under the same *effective context cap C* (Sec. 3): the non-streaming baseline uses a sliding window of
100 length C and recomputes attention each step (no KV cache), while StreamingLLM variants reuse KV
101 cache and apply pruning/re-alignment to stay within the same cap. All methods share the same model,
102 dataset segment, and evaluation protocol implemented in `experiments/eval_streaming_llm.py`.

103 **Protocol and repeatability.** To reduce measurement noise, our reproduction script runs a warmup
104 pass followed by N repeated trials for each configuration and reports mean and standard deviation
105 (saved in JSON). Each JSON record includes environment fingerprints (GPU, driver, torch/CUDA) to
106 prevent reusing baselines across incompatible environments.

107 4.2 Main Results and Ablations

108 We provide a one-click script (`run_paper_experiments.sh`) that runs all configurations and
109 generates LaTeX tables automatically. If the tables are not generated yet, we show placeholders.

110 **Implementation-path fairness check.** To ensure that improvements are not caused by divergent
111 code paths, we also report an *Ours-framework-only* configuration that uses our wrapper implementa-
112 tion but matches Start+Recent semantics (MIT-style cache backend) with all proposed mechanisms
113 disabled.

114 Ablations.

115 5 Discussion

116 **Why do popular inference accelerations fail in batch-1 streaming?** In our setting, streaming
117 bounds attention length, shifting the bottleneck toward MLP and framework/launch overhead. There-
118 fore, further optimizing attention kernels can have limited impact. Moreover, methods that assume
119 static shapes or stable cache semantics may be incompatible with pruning and RoPE re-alignment in
120 streaming decode.

121 **Attempts and analysis of ineffective methods.** We tested several common acceleration routes (see
122 our project log in `docs/`): (i) FlashAttention integration was non-trivial under RoPE re-alignment
123 and did not improve batch-1 TPOT in our environment; (ii) speculative decoding was unreliable
124 on long-form generation and incompatible with streaming cache management; (iii) quantization
125 via TorchAO exhibited numerical issues (INT8 v1) and remained slower even when stabilized

Table 3: Ablation study for Slack and Max_Drop (PG19).

Setting	TPOT \downarrow	Speedup \uparrow	PPL \downarrow
w/o Slack ($\sigma = 0$)	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]
w/o Max_Drop ($\delta = 0$)	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]
Full (Lazy+Slack+Max_Drop)	[INSERT DATA]	[INSERT DATA]	[INSERT DATA]

Table 4: Summary of investigated but ineffective optimization routes in our batch-1 streaming setting.

Method	Outcome	Notes (see project logs)
FlashAttention / FlashDecoding	No gain / hard to integrate	Streaming already bounds attention length; remaining bottlenecks are MLP and launch/framework overhead; RoPE re-alignment complicates clean integration.
Speculative decoding	No gain / unreliable	Long-form generation yields low acceptance; cache consistency with pruning is fragile.
Quantization (TorchAO INT8/INT4)	Slower / unstable	INT8 WO v1 produced NaNs; v2 is stable but slower for batch-1 decode in our stack; INT4 backend dependencies were problematic.
<code>torch.compile</code> / CUDA Graphs	Unstable	Repeated-run CUDA graph overwrite errors observed in rotary-embedding path; shape/cache semantics hinder capture.
HF StaticCache	Incompatible	StaticCache assumes fixed cache updates; pruning can trigger device-side asserts (index out of bounds).
CUDA fusion (residual/LN)	No gain	Amdahl's law: residual/LN is a small fraction; custom kernel launch overhead dominated, leading to slowdown.

(INT8 v2) in batch-1 streaming decode; (iv) `torch.compile` with CUDA Graphs was unstable due to overwritten outputs in rotary-embedding paths; (v) HuggingFace StaticCache conflicted with pruning updates and triggered device-side asserts. These results highlight that strong batch-inference accelerations do not directly transfer to batch-1 streaming decode.

6 Conclusion

We presented three modular KV-cache management mechanisms for StreamingLLM: Lazy Pruning, Slack, and Max_Drop. Our method is training-free and targets long-context decoding where pruning overhead and eviction-induced quality loss become important. We provide a reproducible pipeline and report negative results to clarify which acceleration techniques do or do not apply in this regime.

References

- [1] Charlie Chen, Sebastian Borgeaud, Geoffrey Irving, Jean-Baptiste Lespiau, Laurent Sifre, and John Jumper. Accelerating large language model decoding with speculative sampling. *arXiv preprint arXiv:2302.01318*, 2023.
- [2] Tri Dao. Flashattention-2: Faster attention with better parallelism and work partitioning. *arXiv preprint arXiv:2307.08691*, 2023.
- [3] Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. In *NeurIPS*, 2022.
- [4] Tim Dettmers, Mike Lewis, Younes Belkada, and Luke Zettlemoyer. Llm.int8(): 8-bit matrix multiplication for transformers at scale. *NeurIPS*, 2022.
- [5] Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. Gptq: Accurate post-training quantization for generative pre-trained transformers. In *ICLR*, 2023.

- 147 [6] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu,
148 Joseph E Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large
149 language model serving with pagedattention. *SOSP*, 2023.
- 150 [7] Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast inference from transformers via
151 speculative decoding. In *ICML*, 2023.
- 152 [8] Zichang Liu, Aditya Desai, Fangshuo Liao, Weitao Wang, Victor Xie, Zhaozhuo Xu, Anastasios
153 Kyriolidis, and Anshumali Shrivastava. Scissorhands: Exploiting the persistence of importance
154 hypothesis for llm kv cache compression at test time. *arXiv preprint arXiv:2305.17118*, 2023.
- 155 [9] Jack W Rae, Anna Potapenko, Siddhant M Jayakumar, Chloe Hillier, and Timothy P Lillicrap.
156 Pg-19 language modeling benchmark. *arXiv preprint arXiv:1911.05507*, 2019.
- 157 [10] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient streaming
158 language models with attention sinks. In *ICLR*, 2024.
- 159 [11] Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao
160 Song, Yuandong Tian, Christopher Ré, Clark Barrett, et al. H2o: Heavy-hitter oracle for efficient
161 generative inference of large language models. In *NeurIPS*, 2023.